

ANT208

# What's new in Amazon Athena

Scott Rigney

Sr. Product Manager  
Amazon Web Services

Ofer Eliassaf

Director, REM Cloud Infrastructure  
Mobileye Vision Technologies

# ANT208

## What's new in Amazon Athena

### Speakers

- Ofer Eliassaf, Director, REM Cloud Infrastructure, Mobileye
- Scott Rigney, Sr. Product Manager, Amazon Web Services

### Announcements

- ❖ Lots of incremental announcement. See notes and slide 12 for a summary.

### Takeaways

- ❖ Athena get's faster and new capabilities with SQL v3 engine and Iceberg.

### Azure

- ❖ I don't believe Azure has Athena like capability. Interestingly Azure Data Lake Storage Gen 2 is a new Athena Connector.

[< Section Header](#)

[< Table of Contents](#)

## New Features and Capabilities

- [New](#) – Athena for Apache Spark.
- [New](#) – Athena SQL Engine v3
  - 50+ New SQL Functions, 20% Faster Queries, 100% backwards compatible
  - Faster Queries, on more information, for lower cost
- [New](#) – Faster queries with result reuse

## How MobileEye Use Athena

- MobileEye is an Autonomous Driving software company. They give an [overview](#) of how they chose Athena, and how they use it.

## New Data Lake, Connector, and Integration Features

- Presented AWS's view of Modern Data Architecture
  - Data Lakes – Store all your data in open formats. Cost effective storage, Decouple Storage from Compute, Choice of Analytical and ML engines, process data in place.
- [New](#) – Capabilities for Apache Iceberg (underlying engine)
  - Create Table as Select (allows you to create new tables as a result of a Select query), Views to hide complex logic, Performance 10x faster queries
  - Merge to sync tables and simplify data prep, Vacuum to manage storage, New formats Avro & ORC.
- [New](#) – Expanded fine-grained access controls
- [New](#) – Data source connectors

NOW AVAILABLE

# Amazon Athena for Apache Spark

Get started with interactive analytics on  
Apache Spark in under a second

Harness Apache Spark for  
complex, powerful analytics

Spend more time on insights  
instead of waiting for results

Build applications without managing  
resources or configuring software

GENERALLY AVAILABLE TODAY



© 2022, Amazon Web Services, Inc. or its affiliates. All rights reserved.

NOW AVAILABLE

# Amazon Athena SQL engine v3

Upgrade today to experience multiple  
improvements and new features

Faster queries at lower cost

More frequent updates and  
features

Based on open-source Trino,  
Presto, and AWS-developed  
innovations for both



© 2022, Amazon Web Services, Inc. or its affiliates. All rights reserved.

## SQL engine v3 new features and enhancements

FASTER QUERIES AND MORE FREQUENT UPDATES

NEW



Query result  
caching

NEW



Aggregation  
functions

NEW



Date & time  
functions

UPDATED



Prepared  
statements

NEW



Number  
formatting

UPDATED



Memory  
usage

NEW



More frequent  
updates

UPDATED



Geospatial  
functions

UPDATED



Syntax  
enhancements

UPDATED



AWS Glue  
metadata

NEW



Array  
functions

UPDATED



Right-join  
dynamic filter

NOW AVAILABLE

## Faster Queries with Result Reuse

New caching feature to accelerate repeat  
queries and reduce time to insight

5x

Up to 5x faster queries  
when result reuse is active

\$0

No data scanned, so  
queries are free to you



© 2022, Amazon Web Services, Inc. or its affiliates. All rights reserved.



# The process that led us to Amazon Athena

## Design phase

Initial idea of developing two systems for short/long term storage (cost vs. efficiency)

Started with the long-term storage

Research conclusion: Amazon Athena + S3 parquet files for long term

Solution verification with AWS solution architects

Data-lab in Seattle

## Implementation phase

Conclusion verification and POC implementation in the data-lab

Finalizing implementation

## Production phase

Usage grew exponentially (usage pattern, amount of data, cost)

Iterative optimizations

**Up until now, no need for the short-term storage**



© 2022, Amazon Web Services, Inc. or its affiliates. All rights reserved.

# Optimizations toward our current state

## Iterative approach

- Query-level optimizations
- Partnered with AWS to reach massive concurrency
- Refactored data model to eliminate joins and align with Athena best practices
- Injected partition projection (geospatial + time)

## Our latest optimizations led us to get:

**90%**

Query time decrease

**90%**

Cost decrease

**0**

bottlenecks across different map activities

**10s → 100s**

concurrent map cell runs per minute


**0**

concurrency issues



© 2022, Amazon Web Services, Inc. or its affiliates. All rights reserved.

# Modern data architecture on AWS



© 2022, Amazon Web Services, Inc. or its affiliates. All rights reserved.

# New features and capabilities with Iceberg

NEW



**CTAS**

Easy and fast way to create new tables

NEW



**Views**

Create views to hide complex logic

NEW



**Performance**

10x faster queries when using SQL engine v3



© 2022, Amazon Web Services, Inc. or its affiliates. All rights reserved.



# New features and capabilities with Iceberg

NEW



## Merge

Sync tables to simplify data prep and updates

NEW



## Vacuum

Manage storage and delete records to meet regulatory requirements

NEW



## Avro & ORC

New support for Avro and ORC formats



© 2022, Amazon Web Services, Inc. or its affiliates. All rights reserved.

NOW AVAILABLE

# Expanded fine-grained access controls

Apply AWS Lake Formation access control policies to more of your data

Centrally defined table-, row-, and column-level policies in Lake Formation

Now supports all table formats and file types available with Athena

Filtering operations performed natively by Athena for increased performance



© 2022, Amazon Web Services, Inc. or its affiliates. All rights reserved.

# New data source connectors

DATA SOURCE CONNECTORS FOR IN-PLACE ANALYTICS

## Easy to set up and query

Use the Athena console to create and configure your connectors, then run SQL queries

## Secure cross-account access

Deploy and configure one connector, then use flexible IAM policies to enable cross-account access to other analysts

## Open-source and fully supported

Athena connectors are open-source, so you can use our code as a starting point to building your own

NEW

Azure Data Lake Storage Gen2  
Azure Synapse  
Cloudera  
Google BigQuery  
Hortonworks  
IBM Db2  
Microsoft SQL Server  
Oracle  
SAP HANA  
Snowflake  
Teradata



© 2022, Amazon Web Services, Inc. or its affiliates. All rights reserved.



# What's new with Amazon Athena

KEY TAKEAWAYS TO APPLY TO YOUR WORKLOADS TODAY

- **Amazon Athena for Apache Spark** – Available today! Sub-second startup for interactive Spark analytics. Check out **ANT209** to learn more
- Upgrade to **SQL engine v3** for faster queries and more analytics features
- Turn on **Query Result Reuse** for cached results
- Expanded support of **Apache Iceberg** and **AWS Lake Formation**
- **New data source connectors** and **Amazon AppFlow** help you expand analysis and break down data silos
- And many more ...



© 2022, Amazon Web Services, Inc. or its affiliates. All rights reserved.