

ANT207

# What's new in AWS streaming

Anusha Dasarakothapalli (she/her)

Principal Software Engineer, AWS  
Messaging and Streaming

Divya Sridhar (she/her)

Sr. Product Manager, Amazon Managed  
Streaming for Apache Kafka

# ANT207

## What's new in AWS Streaming

### Speakers

- Divya Sridhar, Senior Product Manager Technical, Amazon
- Anusha Dasarakothapalli, Principal Software Engineer, Amazon

### Announcements

- ❖ Lots of incremental announcement. See the What's New Section and 4<sup>th</sup> Slide in Summary.

### Takeaways

- ❖ Focus on Offloading Capacity Management and Ease of Use.

### Azure

- ❖ Streaming services on Azure are primarily categorized under the [Messaging Services](#) category.

[< Section Header](#)  
[< Table of Contents](#)

## Why Streaming

- Data is a perishable resource, streaming data real-time can help you identify problems before they occur.
- Common streaming-data use cases
  - Anomaly and Fraud Detection
  - Tailoring customer experience in real-time
  - Empowering IoT Analytics
  - Nourishing marketing campaigns
  - Real-time personalization
  - Supporting Healthcare and Emergency Services

## Streaming Portfolio

- AWS Data Streaming, Processing and Analytics
  - Kinesis Data Streams
  - Kinesis Video Streams
  - Kinesis Data Firehose
  - Kinesis Data Analytics
  - Managed streaming for Apache Kafka

## What's New

- Hands Off Capacity Management w/MSK Serverless
  - [New](#) – Use MSK Connect with MSK
  - New – MSK Serverless expanded Region Capability
  - [New](#) – Hands Off Capacity for MSK Tiered Storage
  - [New](#) – Additional Metrics for Kinesis Data Analytics for Apache Flink
- Ease of Use
  - [New](#) – Kinesis Data Streams Viewer in AWS Console
  - [New](#) – Improved Visibility in Kinesis Data Analytics for Apache Flink with Customized Mtx Windows & 5x Performance Increase in Logging
  - [New](#) – MSK Connect now supports private DNS
  - [New](#) – Move data directly from Kinesis and MSK to Redshift

- [New](#) – Kinesis Data Analytics launches support for Apache Flink 1.15.2

# Data streaming workflow



© 2022, Amazon Web Services, Inc. or its affiliates. All rights reserved.

# Common streaming-data use cases



Anomaly and fraud detection



Tailoring customer experience in real time



Empowering IoT analytics



Nourishing marketing campaigns



Real-time personalization



Supporting healthcare and emergency services



© 2022, Amazon Web Services, Inc. or its affiliates. All rights reserved.

# Real-time data streaming, processing, and analytics on AWS

Amazon Kinesis Data Streams



Collect and store data streams for analytics

Amazon Kinesis Video Streams



Collect and store video streams for analytics

Amazon Kinesis Data Firehose



Load data streams into AWS date stores

Amazon Kinesis Data Analytics



Analyze data streams with Apache Flink

Amazon Managed Streaming for Apache Kafka



Fully managed, highly available and secure



© 2022, Amazon Web Services, Inc. or its affiliates. All rights reserved.

# What's new in AWS Streaming

Managed Streaming for Apache Kafka



- ✓ Support for tiered storage
- ✓ MSK Serverless support for MSK Connect
- ✓ MSK Serverless available in 11 AWS Regions
- ✓ MSK Connect support for private DNS
- ✓ Native integration with Amazon Redshift

Kinesis Data Streams



- ✓ Kinesis Data Streams Data Viewer in AWS Management Console
- ✓ Native integration with Amazon Redshift

Kinesis Data Firehose



- ✓ Kinesis Data Firehose supports OpenSearch Serverless as a destination
- ✓ Kinesis Data Firehose supports VPC Flow Logs as a source

Kinesis Data Analytics



- ✓ Enhanced metrics
- ✓ Custom Maintenance Windows
- ✓ Improved application performance with logging
- ✓ Support for Apache Flink 1.15.2




© 2022, Amazon Web Services, Inc. or its affiliates. All rights reserved.



# Amazon MSK Serverless

EASILY STREAM DATA WITH AMAZON MSK WITHOUT MANAGING CLUSTER CAPACITY

"Our banking application gets used mostly during business hours. We have to scale our cluster up and down to stay cost-efficient."



"We have sales days every month, and we have to keep rescaling our clusters based on demand. Having an autoscaling feature will reduce our operational overhead."

Customer



On-demand streaming capacity



Fully compatible



Throughput based pricing



Same security as Amazon MSK



Auto partition management



Same high availability



- Use Amazon MSK Connect with Amazon MSK
- MSK Serverless expanded Region availability



© 2022, Amazon Web Services, Inc. or its affiliates. All rights reserved.

# Amazon MSK support for tiered storage



VIRTUALLY UNLIMITED AND COST-EFFECTIVE STORAGE TIER FOR LONG-LIVED DATA IN APACHE KAFKA APPLICATIONS



Tiered storage provides a new low-cost option that scales to virtually unlimited storage

- Increase your topics' storage without needing to increase the number of brokers
- Reprocess old data in the exact same order in which it was produced
- Save with a low-cost storage tier



© 2022, Amazon Web Services, Inc. or its affiliates. All rights reserved.


# Kinesis Data Analytics for Apache Flink enhanced metrics



Kinesis Data Analytics for Apache Flink task managers

Amazon CloudWatch

Use metrics to scale Kinesis Data Analytics Apache Flink applications



© 2022, Amazon Web Services, Inc. or its affiliates. All rights reserved.

# Kinesis Data Streams data viewer in AWS Console



INTERACTIVE, NO-CODE DATA EXPLORER

Use cases

1. Preview data quickly
2. Understand record schema
3. Data debugging

Name	Status	Capacity mode	Provisioned shards	Data retention period	Encryption	Consumers with enhanced fan-out
001-001-DELETE	Active	Provisioned	1	1 day	Enabled	0
001-001-DELETE	Active	Provisioned	1	1 day	Disabled	0
Example-1	Active	Provisioned	4	30 days	Disabled	0
Example-2	Active	Provisioned	8	30 days	Disabled	0
Example-3	Active	Provisioned	8	30 days	Disabled	0
Example-4	Active	Provisioned	4	30 days	Disabled	0
Example-5	Active	Provisioned	1	30 days	Enabled	0
Example-6	Active	On-demand	-	30 days	Enabled	0
Example-7	Active	On-demand	-	1 day	Disabled	0
Example-8	Active	On-demand	-	1 day	Disabled	0
Example-9	Active	Provisioned	4	1 day	Disabled	0
Example-10	Active	Provisioned	4	1 day	Disabled	0
Example-11	Active	Provisioned	1	1 day	Disabled	0
Example-12	Active	Provisioned	1	1 day	Disabled	0
Example-13	Active	On-demand	-	1 day	Enabled	0
Example-14	Active	Provisioned	4	1 day	Disabled	0
Example-15	Active	Provisioned	10	1 day	Enabled	1



© 2022, Amazon Web Services, Inc. or its affiliates. All rights reserved.

# Kinesis Data Analytics for Apache Flink: Improved usability



Customize maintenance windows



5x improvement in application performance with logging turned on



© 2022, Amazon Web Services, Inc. or its affiliates. All rights reserved.

# Amazon MSK Connect

EASILY MOVE DATA BETWEEN APACHE KAFKA CLUSTERS AND OTHER SYSTEMS



Fully managed



Fully compatible



Serverless



Autoscaling



**Amazon MSK Connect private DNS support**

- ✓ Keep databases, data warehouses, and other systems like the secrets manager in your own VPC, inaccessible to the internet
- ✓ Use Amazon MSK Connect with Amazon MSK Serverless



© 2022, Amazon Web Services, Inc. or its affiliates. All rights reserved.

# Move data from MSK and Kinesis directly to Amazon Redshift



NATIVE INTEGRATION WITH STREAMING ENGINES FOR FAST AND EASY DATA INGESTION

Create materialized views on top of streams directly


Achieve low latency in the seconds

Support up to hundreds of MBps of throughput

Entirely configurable with SQL

Define a schema or ingest semi-structured data with SUPER & VARBYTE data types

Learn more: <https://go.aws/3eG5kZF>



© 2022, Amazon Web Services, Inc. or its affiliates. All rights reserved.

# Amazon Kinesis Data Analytics launches support for Apache Flink version 1.15.2

Kinesis Data Analytics championed Apache Flink Version 1.15.2 release



Focus feature: Async Sink, an AWS contributed feature



A new framework allowing customers to easily create custom sinks



AWS facilitating community project to create a DynamoDB sink using the framework

Build a custom sink in under 20 minutes: <https://bit.ly/3VyAkeJ>



© 2022, Amazon Web Services, Inc. or its affiliates. All rights reserved.